

# HX Infrastructure Inventory Documentation

---



## Comprehensive Inventory Structure and Visual Diagrams

---

This document provides detailed visual documentation for the HX Infrastructure inventory structure, showing environment organization, host groupings, service relationships, and variable inheritance patterns.

## Environment Overview Diagram

```
graph TB
    subgraph "HX Infrastructure Environments"
        subgraph "Development Environment"
            DEV_DOMAIN[Domain: dev-test.hana-x.ai<br/>Network: 192.168.10.0/24<br/>15 Servers]

            subgraph "Dev Services"
                DEV_INFRA[Infrastructure<br/>🏢 DC + CA<br/>2 servers]
                DEV_AI[AI/ML Services<br/>🤖 LiteLLM + Ollama<br/>4 servers]
                DEV_UI[UI Services<br/>🌐 WebUI + LB<br/>3 servers]
                DEV_OPS[Operations<br/>🔧 DB + Cache + Monitor<br/>6 servers]
            end
        end

        subgraph "Test Environment"
            TEST_DOMAIN[Domain: test.hana-x.ai<br/>Network: 192.168.20.0/24<br/>8 Servers]

            subgraph "Test Services"
                TEST_INFRA[Infrastructure<br/>🏢 DC Only<br/>1 server]
                TEST_AI[AI/ML Services<br/>🤖 CPU-only Testing<br/>2 servers]
                TEST_UI[UI Services<br/>🌐 Basic WebUI<br/>2 servers]
                TEST_OPS[Operations<br/>🔧 Essential Services<br/>3 servers]
            end
        end

        subgraph "Production Environment"
            PROD_DOMAIN[Domain: hana-x.ai<br/>Network: 10.0.0.0/16<br/>22 Servers]

            subgraph "Prod Services"
                PROD_INFRA[Infrastructure<br/>🏢 HA DC + CA<br/>4 servers]
                PROD_AI[AI/ML Services<br/>🤖 GPU Cluster<br/>6 servers]
                PROD_UI[UI Services<br/>🌐 HA WebUI + LB<br/>5 servers]
                PROD_OPS[Operations<br/>🔧 Enterprise Grade<br/>7 servers]
            end
        end

        DEV_INFRA --> DEV_AI
        DEV_AI --> DEV_UI
        DEV_UI --> DEV_OPS

        TEST_INFRA --> TEST_AI
        TEST_AI --> TEST_UI
        TEST_UI --> TEST_OPS

        PROD_INFRA --> PROD_AI
        PROD_AI --> PROD_UI
        PROD_UI --> PROD_OPS

        %% Styling
        classDef devClass fill:#e3f2fd,stroke:#1976d2,stroke-width:2px
        classDef testClass fill:#f3e5f5,stroke:#7b1fa2,stroke-width:2px
        classDef prodClass fill:#e8f5e8,stroke:#388e3c,stroke-width:2px

        class DEV_DOMAIN,DEV_INFRA,DEV_AI,DEV_UI,DEV_OPS devClass
        class TEST_DOMAIN,TEST_INFRA,TEST_AI,TEST_UI,TEST_OPS testClass
        class PROD_DOMAIN,PROD_INFRA,PROD_AI,PROD_UI,PROD_OPS prodClass
```

## **Service Grouping Structure**

---

```

graph TB
  subgraph "HX Infrastructure Service Groups"
    subgraph "Infrastructure Services"
      DC[Domain Controllers<br/>🏢 Active Directory<br/>DNS, DHCP, GPO<br/>Windows Server]
      CA[Certificate Authorities<br/>🔒 PKI Infrastructure<br/>SSL Certificates<br/>Code Signing]
    end

    subgraph "AI/ML Services"
      LITELLM[LiteLLM Gateway<br/>🏠 API Gateway<br/>Model Routing<br/>Rate Limiting]
      OLLAMA[Ollama Inference<br/>🧠 Local LLMs<br/>GPU Acceleration<br/>Model Management]
      MODELS[Model Storage<br/>📦 Model Repository<br/>Version Control<br/>Registry Service]
    end

    subgraph "UI Services"
      WEBUI[Web Interface<br/>💬 Chat Interface<br/>User Management<br/>Admin Panel]
      LB[Load Balancers<br/>⚖️ Traffic Distribution<br/>SSL Termination<br/>High Availability]
    end

    subgraph "Operations Services"
      DB[PostgreSQL Cluster<br/>🗄️ Primary Database<br/>Replication<br/>Backup & Recovery]
      CACHE[Redis Cluster<br/>⚡ Caching Layer<br/>Session Storage<br/>Pub/Sub Messaging]
      MONITOR[Monitoring Stack<br/>📊 Prometheus + Grafana<br/>Metrics Collection<br/>Alerting]
      LOGS[Logging Stack<br/>📝 ELK Stack<br/>Log Aggregation<br/>Search & Analysis]
    end
  end

  %% Service Dependencies
  DC --> CA
  CA --> LB
  LB --> WEBUI
  WEBUI --> LITELLM
  LITELLM --> OLLAMA
  OLLAMA --> MODELS
  WEBUI --> DB
  WEBUI --> CACHE

  %% Monitoring Connections
  MONITOR -. -> DC
  MONITOR -. -> CA
  MONITOR -. -> LITELLM
  MONITOR -. -> OLLAMA
  MONITOR -. -> WEBUI
  MONITOR -. -> LB
  MONITOR -. -> DB
  MONITOR -. -> CACHE

  %% Logging Connections
  LOGS -. -> DC
  LOGS -. -> CA
  LOGS -. -> LITELLM
  LOGS -. -> OLLAMA

```

```
LOGS --> WEBUI
LOGS --> LB
LOGS --> DB
LOGS --> CACHE
LOGS --> MONITOR

%% Styling
classDef infraClass fill:#ffebee,stroke:#d32f2f,stroke-width:2px
classDef aiClass fill:#e8f5e8,stroke:#388e3c,stroke-width:2px
classDef uiClass fill:#e3f2fd,stroke:#1976d2,stroke-width:2px
classDef opsClass fill:#fff3e0,stroke:#f57c00,stroke-width:2px

class DC,CA infraClass
class LITELLM,OLLAMA,MODELS aiClass
class WEBUI,LB uiClass
class DB,CACHE,MONITOR,LOGS opsClass
```

## **Network Topology by Environment**

---

### **Development Environment Network**

```

graph TB
  subgraph "Development Network - 192.168.10.0/24"
    subgraph "Infrastructure Subnet - .10-.19"
      DC_DEV[hx-dc-01<br/>192.168.10.10<br/>Domain Controller]
      CA_DEV[hx-ca-01<br/>192.168.10.11<br/>Certificate Authority]
    end

    subgraph "AI/ML Subnet - .20-.29"
      LITELLM_DEV[hx-litellm-01<br/>192.168.10.20<br/>LiteLLM Gateway]
      OLLAMA1_DEV[hx-ollama-01<br/>192.168.10.21<br/>Ollama Primary]
      OLLAMA2_DEV[hx-ollama-02<br/>192.168.10.22<br/>Ollama Secondary]
      MODELS_DEV[hx-models-01<br/>192.168.10.23<br/>Model Storage]
    end

    subgraph "UI Subnet - .30-.39"
      WEBUI1_DEV[hx-webui-01<br/>192.168.10.30<br/>WebUI Primary]
      WEBUI2_DEV[hx-webui-02<br/>192.168.10.31<br/>WebUI Secondary]
      LB_DEV[hx-lb-01<br/>192.168.10.32<br/>Load Balancer<br/>VIP: .100]
    end

    subgraph "Operations Subnet - .40-.59"
      DB1_DEV[hx-db-01<br/>192.168.10.40<br/>PostgreSQL Master]
      DB2_DEV[hx-db-02<br/>192.168.10.41<br/>PostgreSQL Replica]
      REDIS1_DEV[hx-redis-01<br/>192.168.10.42<br/>Redis Master]
      REDIS2_DEV[hx-redis-02<br/>192.168.10.43<br/>Redis Replica]
      MON_DEV[hx-monitor-01<br/>192.168.10.50<br/>Monitoring]
      LOG_DEV[hx-logs-01<br/>192.168.10.51<br/>Logging]
    end

    GATEWAY_DEV[Gateway<br/>192.168.10.1]
  end

  %% Network Connections
  GATEWAY_DEV --> LB_DEV
  LB_DEV --> WEBUI1_DEV
  LB_DEV --> WEBUI2_DEV
  WEBUI1_DEV --> LITELLM_DEV
  WEBUI2_DEV --> LITELLM_DEV
  LITELLM_DEV --> OLLAMA1_DEV
  LITELLM_DEV --> OLLAMA2_DEV
  OLLAMA1_DEV --> MODELS_DEV
  OLLAMA2_DEV --> MODELS_DEV
  WEBUI1_DEV --> DB1_DEV
  WEBUI2_DEV --> DB1_DEV
  DB1_DEV --> DB2_DEV
  WEBUI1_DEV --> REDIS1_DEV
  WEBUI2_DEV --> REDIS1_DEV
  REDIS1_DEV --> REDIS2_DEV

  %% Styling
  classDef infraClass fill:#ffebee,stroke:#d32f2f,stroke-width:2px
  classDef aiClass fill:#e8f5e8,stroke:#388e3c,stroke-width:2px
  classDef uiClass fill:#e3f2fd,stroke:#1976d2,stroke-width:2px
  classDef opsClass fill:#fff3e0,stroke:#f57c00,stroke-width:2px
  classDef netClass fill:#f3e5f5,stroke:#7b1fa2,stroke-width:2px

  class DC_DEV,CA_DEV infraClass
  class LITELLM_DEV,OLLAMA1_DEV,OLLAMA2_DEV,MODELS_DEV aiClass
  class WEBUI1_DEV,WEBUI2_DEV,LB_DEV uiClass
  class DB1_DEV,DB2_DEV,REDIS1_DEV,REDIS2_DEV,MON_DEV,LOG_DEV opsClass
  class GATEWAY_DEV netClass

```

**Production Environment Network**



```

graph TB
  subgraph "Production Network - 10.0.0.0/16"
    subgraph "Infrastructure Tier - 10.0.1.0/24"
      subgraph "AZ-1"
        DC1_PROD[hx-dc-prod-01<br/>10.0.1.10<br/>Primary DC]
        CA1_PROD[hx-ca-prod-01<br/>10.0.1.20<br/>Primary CA]
      end
      subgraph "AZ-2"
        DC2_PROD[hx-dc-prod-02<br/>10.0.1.11<br/>Secondary DC]
        CA2_PROD[hx-ca-prod-02<br/>10.0.1.21<br/>Secondary CA]
      end
    end

    subgraph "AI/ML Tier - 10.0.2.0/24"
      subgraph "AZ-1"
        LITELLM1_PROD[hx-litellm-prod-01<br/>10.0.2.10<br/>LiteLLM Primary]
        OLLAMA1_PROD[hx-ollama-prod-01<br/>10.0.2.20<br/>Ollama Primary]
        OLLAMA3_PROD[hx-ollama-prod-03<br/>10.0.2.22<br/>Ollama Tertiary]
        MODELS1_PROD[hx-models-prod-01<br/>10.0.2.30<br/>Model Primary]
      end
      subgraph "AZ-2"
        LITELLM2_PROD[hx-litellm-prod-02<br/>10.0.2.11<br/>LiteLLM Secondary]
        OLLAMA2_PROD[hx-ollama-prod-02<br/>10.0.2.21<br/>Ollama Secondary]
        MODELS2_PROD[hx-models-prod-02<br/>10.0.2.31<br/>Model Replica]
      end
    end

    subgraph "UI Tier - 10.0.3.0/24"
      subgraph "AZ-1"
        WEBUI1_PROD[hx-webui-prod-01<br/>10.0.3.10<br/>WebUI Primary]
        WEBUI3_PROD[hx-webui-prod-03<br/>10.0.3.12<br/>WebUI Tertiary]
        LB1_PROD[hx-lb-prod-01<br/>10.0.3.20<br/>LB Primary<br/>VIP: .100]
      end
      subgraph "AZ-2"
        WEBUI2_PROD[hx-webui-prod-02<br/>10.0.3.11<br/>WebUI Secondary]
        LB2_PROD[hx-lb-prod-02<br/>10.0.3.21<br/>LB Secondary<br/>VIP: .100]
      end
    end

    subgraph "Data Tier - 10.0.4.0/24"
      subgraph "AZ-1"
        DB1_PROD[hx-db-prod-01<br/>10.0.4.10<br/>PostgreSQL Master]
        DB3_PROD[hx-db-prod-03<br/>10.0.4.12<br/>PostgreSQL Replica 2]
        REDIS1_PROD[hx-redis-prod-01<br/>10.0.4.20<br/>Redis Master]
        REDIS3_PROD[hx-redis-prod-03<br/>10.0.4.22<br/>Redis Replica 2]
      end
      subgraph "AZ-2"
        DB2_PROD[hx-db-prod-02<br/>10.0.4.11<br/>PostgreSQL Replica 1]
        REDIS2_PROD[hx-redis-prod-02<br/>10.0.4.21<br/>Redis Replica 1]
      end
    end

    subgraph "Management Tier - 10.0.5.0/24"
      subgraph "AZ-1"
        MON1_PROD[hx-monitor-prod-01<br/>10.0.5.10<br/>Monitoring Primary]
        LOG1_PROD[hx-logs-prod-01<br/>10.0.5.20<br/>Logging Primary]
      end
      subgraph "AZ-2"
        MON2_PROD[hx-monitor-prod-02<br/>10.0.5.11<br/>Monitoring Secondary]
        LOG2_PROD[hx-logs-prod-02<br/>10.0.5.21<br/>Logging Secondary]
      end
    end
  end

```

```

end

%% High Availability Connections
DC1_PROD <--> DC2_PROD
CA1_PROD <--> CA2_PROD
LITELLM1_PROD <--> LITELLM2_PROD
OLLAMA1_PROD <--> OLLAMA2_PROD
MODELS1_PROD <--> MODELS2_PROD
LB1_PROD <--> LB2_PROD
DB1_PROD --> DB2_PROD
DB1_PROD --> DB3_PROD
REDIS1_PROD --> REDIS2_PROD
REDIS1_PROD --> REDIS3_PROD
MON1_PROD <--> MON2_PROD
LOG1_PROD <--> LOG2_PROD

%% Styling
classDef az1Class fill:#e3f2fd,stroke:#1976d2,stroke-width:2px
classDef az2Class fill:#f3e5f5,stroke:#7b1fa2,stroke-width:2px

class DC1_PROD,CA1_PROD,LITELLM1_PROD,OLLAMA1_PROD,OLLAMA3_PROD,MODELS1_PROD,WEBUI1_PROD,WEBUI3_PROD,LB1_PROD,DB1_PROD,DB3_PROD,REDIS1_PROD,REDIS3_PROD,MON1_PROD,LOG1_PROD az1Class
class DC2_PROD,CA2_PROD,LITELLM2_PROD,OLLAMA2_PROD,MODELS2_PROD,WEBUI2_PROD,LB2_PROD,DB2_PROD,REDIS2_PROD,MON2_PROD,LOG2_PROD az2Class

```



## Variable Inheritance Hierarchy

---

```

graph TB
    subgraph "Ansible Variable Precedence (High to Low)"
        EXTRA[Extra Variables<br/>🔓 ansible-playbook -e<br/>Highest Priority<br/>Runtime Overrides]

        HOST_VARS[Host Variables<br/>📁 host_vars/<hostname>.yaml<br/>Host-specific Settings<br/>Individual Customization]

        subgraph "Group Variables"
            GROUP_SPECIFIC[Specific Groups<br/>📁 group_vars/infrastructure.yaml<br/>📁 group_vars/ai_ml.yaml<br/>📁 group_vars/ui.yaml<br/>📁 group_vars/operations.yaml]

            GROUP_ALL[All Group<br/>📁 group_vars/all.yaml<br/>Global Defaults<br/>Common Configuration]
        end

        subgraph "Role Variables"
            ROLE_VARS[Role Variables<br/>📁 roles/*/vars/main.yaml<br/>Role-specific Values<br/>High Priority within Role]

            ROLE_DEFAULTS[Role Defaults<br/>📁 roles/*/defaults/main.yaml<br/>Default Values<br/>Lowest Priority]
        end

        subgraph "Inventory Variables"
            INV_HOST[Inventory Host Vars<br/>📁 inventory/host_vars/<hostname>.yaml<br/>>Environment-specific Host Settings]

            INV_GROUP[Inventory Group Vars<br/>📁 inventory/group_vars/<group>.yaml<br/>>Environment-specific Group Settings]
        end

        subgraph "Secrets Management"
            VAULT[Ansible Vault<br/>🔒 Encrypted Variables<br/>Sensitive Data<br/>Cross-referenced in other vars]
        end
    end

    %% Precedence Flow (Higher to Lower)
    EXTRA --> HOST_VARS
    HOST_VARS --> GROUP_SPECIFIC
    GROUP_SPECIFIC --> GROUP_ALL
    GROUP_ALL --> INV_HOST
    INV_HOST --> INV_GROUP
    INV_GROUP --> ROLE_VARS
    ROLE_VARS --> ROLE_DEFAULTS

    %% Vault Integration (dotted lines show cross-references)
    VAULT -. -> EXTRA
    VAULT -. -> HOST_VARS
    VAULT -. -> GROUP_SPECIFIC
    VAULT -. -> GROUP_ALL
    VAULT -. -> INV_HOST
    VAULT -. -> INV_GROUP

    %% Styling
    classDef highPriority fill:#ffccdd,stroke:#d32f2f,stroke-width:3px
    classDef mediumPriority fill:#fff3e0,stroke:#f57c00,stroke-width:2px
    classDef lowPriority fill:#e8f5e8,stroke:#388e3c,stroke-width:2px
    classDef secretClass fill:#f3e5f5,stroke:#7b1fa2,stroke-width:2px

    class EXTRA highPriority

```

```
class HOST_VARS, GROUP_SPECIFIC mediumPriority
class GROUP_ALL, INV_HOST, INV_GROUP, ROLE_VARS mediumPriority
class ROLE_DEFAULTS lowPriority
class VAULT secretClass
```

## Environment-Specific Configuration Patterns

---

```

graph TB
  subgraph "Configuration Management Pattern"
    subgraph "Base Configuration"
      BASE_ALL[group_vars/all.yml<br/>🌐 Global Defaults<br/>Common to all environments<br/>Security, monitoring, backup settings]

      BASE_GROUPS[Group Variables<br/>📁 group_vars/<service>.yaml<br/>Service-specific defaults<br/>PostgreSQL, Redis, Nginx configs]
    end

    subgraph "Environment Overrides"
      ENV_DEV[Development<br/>🔧 inventories/dev/<br/>Debug enabled<br/>Relaxed security<br/>Test data enabled]

      ENV_TEST[Test<br/>🔧 inventories/test/<br/>Validation enabled<br/>Performance testing<br/>Minimal resources]

      ENV_PROD[Production<br/>🔒 inventories/prod/<br/>Security hardened<br/>High availability<br/>Performance optimized]
    end

    subgraph "Host-Specific Customization"
      HOST_CUSTOM[Host Variables<br/>📁 host_vars/<hostname>.yaml<br/>Individual server settings<br/>Resource allocation<br/>Special configurations]
    end

    subgraph "Runtime Flexibility"
      RUNTIME[Runtime Variables<br/>🔧 Extra vars (-e)<br/>Deployment-specific<br/>Feature flags<br/>Emergency overrides]
    end
  end

  %% Configuration Flow
  BASE_ALL --> ENV_DEV
  BASE_ALL --> ENV_TEST
  BASE_ALL --> ENV_PROD

  BASE_GROUPS --> ENV_DEV
  BASE_GROUPS --> ENV_TEST
  BASE_GROUPS --> ENV_PROD

  ENV_DEV --> HOST_CUSTOM
  ENV_TEST --> HOST_CUSTOM
  ENV_PROD --> HOST_CUSTOM

  HOST_CUSTOM --> RUNTIME

  %% Example Overrides
  ENV_DEV -.-> |"debug_mode: true<br/>ssl_verification: false<br/>log_level: DEBUG"| HOST_CUSTOM
  ENV_TEST -.-> |"test_mode: true<br/>performance_testing: true<br/>cleanup_enabled: true"| HOST_CUSTOM
  ENV_PROD -.-> |"security_hardening: true<br/>high_availability: true<br/>monitoring_enhanced: true"| HOST_CUSTOM

  %% Styling
  classDef baseClass fill:#e3f2fd,stroke:#1976d2,stroke-width:2px
  classDef envClass fill:#f3e5f5,stroke:#7b1fa2,stroke-width:2px
  classDef customClass fill:#e8f5e8,stroke:#388e3c,stroke-width:2px
  classDef runtimeClass fill:#ffebee,stroke:#d32f2f,stroke-width:2px

  class BASE_ALL,BASE_GROUPS baseClass

```

```
class ENV_DEV,ENV_TEST,ENV_PROD envClass  
class HOST_CUSTOM customClass  
class RUNTIME runtimeClass
```



## **Deployment Target Groups**

---

```

graph TB
  subgraph "Deployment Target Groupings"
    subgraph "Functional Groups"
      INFRA_GROUP[Infrastructure Services<br/>🏢 Domain Controllers<br/>👤 Certificate Authorities<br/>Windows-based services]

      AI_GROUP[AI/ML Services<br/>🧠 LiteLLM Gateway<br/>🧠 Ollama Inference<br/>📦 Model Storage<br/>GPU-enabled hosts]

      UI_GROUP[UI Services<br/>💬 Web Interface<br/>⚖️ Load Balancers<br/>🌐 Frontend services]

      OPS_GROUP[Operations Services<br/>🗄️ PostgreSQL Cluster<br/>⚡ Redis Cache<br/>📊 Monitoring Stack<br/>📝 Logging Stack]
    end

    subgraph "Platform Groups"
      LINUX_HOSTS[Linux Hosts<br/>🐧 Ubuntu/RHEL<br/>Container-ready<br/>Most AI/ML and web services]

      WINDOWS_HOSTS[Windows Hosts<br/>🖥️ Windows Server<br/>Active Directory<br/>Certificate Services]
    end

    subgraph "Role-Based Groups"
      PRIMARY[Primary Services<br/>🎯 Master nodes<br/>Main service instances<br/>High priority]

      SECONDARY[Secondary Services<br/>🔄 Replica nodes<br/>Backup instances<br/>Failover targets]

      GPU_HOSTS[GPU-Enabled Hosts<br/>🎮 NVIDIA GPUs<br/>CUDA support<br/>AI/ML inference]
    end

    subgraph "Availability Groups"
      AZ1[Availability Zone 1<br/>🏢 Primary datacenter<br/>Main service instances<br/>Active workloads]

      AZ2[Availability Zone 2<br/>🏢 Secondary datacenter<br/>Replica instances<br/>Disaster recovery]
    end

    subgraph "Maintenance Groups"
      CRITICAL[Critical Services<br/>🔥 Zero-downtime required<br/>Database masters<br/>Load balancers]

      STANDARD[Standard Services<br/>⚙️ Planned maintenance windows<br/>Application servers<br/>Monitoring services]
    end
  end

  %% Group Relationships
  INFRA_GROUP --> WINDOWS_HOSTS
  AI_GROUP --> LINUX_HOSTS
  AI_GROUP --> GPU_HOSTS
  UI_GROUP --> LINUX_HOSTS
  OPS_GROUP --> LINUX_HOSTS

  PRIMARY --> CRITICAL
  SECONDARY --> STANDARD

```

```

AZ1 --> PRIMARY
AZ2 --> SECONDARY

%% Styling
classDef functionalClass fill:#e3f2fd,stroke:#1976d2,stroke-width:2px
classDef platformClass fill:#f3e5f5,stroke:#7b1fa2,stroke-width:2px
classDef roleClass fill:#e8f5e8,stroke:#388e3c,stroke-width:2px
classDef availabilityClass fill:#fff3e0,stroke:#f57c00,stroke-width:2px
classDef maintenanceClass fill:#ffebee,stroke:#d32f2f,stroke-width:2px

class INFRA_GROUP,AI_GROUP,UI_GROUP,OPS_GROUP functionalClass
class LINUX_HOSTS,WINDOWS_HOSTS platformClass
class PRIMARY,SECONDARY,GPU_HOSTS roleClass
class AZ1,AZ2 availabilityClass
class CRITICAL,STANDARD maintenanceClass

```

## Inventory Usage Examples

### Basic Deployment Commands

```

# Deploy entire development environment
ansible-playbook -i inventories/dev/hosts.yml playbooks/site/main.yml

# Deploy only AI/ML services in production
ansible-playbook -i inventories/prod/hosts.yml playbooks/services/ai-ml.yml --limit ai_
ml

# Deploy to specific host group
ansible-playbook -i inventories/prod/hosts.yml playbooks/services/database.yml --limit
databases

# Deploy with environment-specific overrides
ansible-playbook -i inventories/prod/hosts.yml playbooks/site/main.yml -e "environ-
ment=production" -e "debug_mode=false"

# Deploy to primary services only
ansible-playbook -i inventories/prod/hosts.yml playbooks/maintenance/update.yml --lim-
it primary_services

# Deploy to specific availability zone
ansible-playbook -i inventories/prod/hosts.yml playbooks/services/monitoring.yml --lim-
it availability_zone_1

```

## Maintenance Operations

```
# Update only GPU-enabled hosts
ansible-playbook -i inventories/prod/hosts.yml playbooks/maintenance/gpu-update.yml --limit gpu_enabled_hosts

# Backup all databases
ansible-playbook -i inventories/prod/hosts.yml playbooks/maintenance/backup.yml --limit databases

# Restart secondary services (safe maintenance)
ansible-playbook -i inventories/prod/hosts.yml playbooks/maintenance/restart.yml --limit secondary_services

# Check health of critical services
ansible-playbook -i inventories/prod/hosts.yml playbooks/monitoring/health-check.yml --limit critical_services
```

## Environment-Specific Operations

```
# Development environment with debug enabled
ansible-playbook -i inventories/dev/hosts.yml playbooks/site/main.yml -e "debug_mode=true" -e "log_level=DEBUG"

# Test environment with performance testing
ansible-playbook -i inventories/test/hosts.yml playbooks/testing/performance.yml -e "load_testing=true"

# Production deployment with security hardening
ansible-playbook -i inventories/prod/hosts.yml playbooks/site/main.yml -e "security_hardening=true" -e "high_availability=true"
```

---

This comprehensive inventory documentation provides clear visual representations of the HX Infrastructure organization, making it easy to understand the relationships between environments, services, and configuration management patterns.